

ADDING SYSTEM CALLS IN XV6 OPERATING SYSTEM

Project Report

Subject: (IT-204) OPERATING SYSTEM DESIGN

**BACHELOR OF TECHNOLOGY IN
INFORMATION TECHNOLOGY**

4th Semester

Submitted by: VARUN KUMAR (2K19/IT/140)

YASHIT KUMAR (2K19/IT/149)

Under the Supervision of

Mr. Jasraj Meena



**DEPARTMENT OF INFORMATION TECHNOLOGY DELHI TECHNOLOGICAL
UNIVERSITY (Formerly Delhi College of Engineering) Bawana Road, Delhi-110042**

DELHI TECHNOLOGICAL UNIVERSITY

(Formerly Delhi College of Engineering)

Bawana Road, Delhi-110042

CANDIDATE'S DECLARATION

We Varun Kumar, Roll No – 2K19/IT/140 & Yashit Kumar, Roll No - 2K19/IT/149, students of B.Tech. (INFORMATION TECHNOLOGY), hereby declare that the project Dissertation titled “**Adding a system call in XV6 OS**” which is submitted by us to the Department of INFORMATION TECHNOLOGY, Delhi Technological University, Delhi in partial fulfillment of the requirement for the award of the 4th semester of the Bachelor of Technology, is made by us. This work has not previously formed the basis for the award of any Degree, Diploma Associateship, Fellowship or other similar title or recognition.

Place: Delhi

Date: 26-05-2021

Varun Kumar

Yashit Kumar

DELHI TECHNOLOGICAL UNIVERSITY

(Formerly Delhi College of Engineering)

Bawana Road, Delhi-110042

CERTIFICATE

I hereby certify that the Project: " **Adding a system call in XV6 OS** " which is submitted by Varun Kumar, Roll No – 2K19/IT/140 & Yashit Kumar, Roll No – 2K19/IT/149, INFORMATION TECHNOLOGY, Delhi Technological University, Delhi in fulfillment of the requirement for the 4th semester of Bachelor of Technology, is record of the project work carried out by the students under my supervision. To the best of my knowledge this work has not been submitted in part or full for any Degree or Diploma to this University or elsewhere.

Place: Delhi

Supervisor : Mr. Jasraj Meena

Date: 26-May-2021

(Assistant Professor)

DEPARTMENT OF INFORMATION TECHNOLOGY
DELHI TECHNOLOGICAL UNIVERSITY

(Formerly Delhi College of Engineering)

Bawana Road, Delhi-110042

ACKNOWLEDGEMENT

We would like to convey our heartfelt thanks to our supervisor Mr. Jasraj Meena for his ingenious ideas, tremendous help and cooperation.

We are extremely grateful to our friends who gave valuable suggestions and guidance for completion of our project. The cooperation and healthy criticism came handy and useful with them.

Finally we would like to thank all the above mentioned people once again.

We would also like to thank our college, Delhi Technological University for giving us an opportunity to study an important subject like **“Operating System”** which would help in our holistic development as an engineer.

DEPARTMENT OF INFORMATION TECHNOLOGY

DELHI TECHNOLOGICAL UNIVERSITY

(Formerly Delhi College of Engineering)

Bawana Road, Delhi-110042

ABSTRACT

The idea of the project is to add system calls to the Linux Kernel xv6 developed by MIT UNIVERSITY. We have used the concepts of operating systems and have included the explanation of process and commands required in kernel compilation, and briefly explained the process of system call addition, commands and packages used to do so. We have tried our best to make the code as efficient as possible.

The objectives of this project were to extend the concepts of Operating System to real world activities. The objectives were fulfilled by learning the concepts of system call, priority scheduler, round robin scheduler and many others throughout the course of four months, simultaneously studying research papers about them and applying them by writing code which improved our skills of fixing bugs, problems and errors and helped in improving our concepts while gaining experience of working in a team

TABLE OF CONTENTS

S.No	Content	Page No.
1.	Objectives and Keywords	7
2.	Introduction	8
3.	System Requirements	9
4.	Methodology	11
5.	Added System Calls	16
6.	Output	20
7.	Conclusion	23
8.	References	24

OBJECTIVES

- To create a system call that can be used to count the total number of method calls in the Linux Kernel xv6 developed by MIT University.
- To create a system call that can be used to change the priority of the system call.
- To create a system call that is used for dummy calculations by creating a parent-child processes.
- To create a system call that shows the state and priority of process.

KEYWORDS

- **System Call**
- **Linux Kernel xv6**
- **QEMU**

INTRODUCTION

A System call is a mechanism that provides the interface between a process and the operating system.

A computer program makes a system call when it makes a request to the operating system's kernel.

System call **provides** the services of the operating system to the user programs via Application Program Interface(API).System calls are the only entry points into the kernel system.

Services Provided by System Calls :

1. Process creation and management
2. Main memory management
3. File Access, Directory and File system management
4. Device handling(I/O)
5. Protection
6. Networking, etc.

SYSTEM REQUIREMENTS

➤ **Operating System=> Linux xv6 developed by MIT University.** ➤ **Virtual Box**

➤ **Quick Emulator (QEMU)**

➤ **Language Used=> C**

PREREQUISITES

There are certain prerequisites for adding a system call in Linux kernel XV6:

- Use a sudo user, so as to avoid unnecessary issues.

- Use the most recent updates of Ubuntu.

To check version:

`lsb_release -a`

To update ubuntu:

- The following packages are required for compilation and is a must. Run commands:

- Git is a version control. It is required to clone repositories. To install:

`sudo apt-get install git`

- Get the latest version of qemu based on computer's hardware.

- Get XV6 OS from official MIT Github link.

LINUX XV6 OS

- xv6 is a modern reimplementation of Sixth Edition Unix in ANSI C for multiprocessor x86 and RISC-V systems
- This is a lightweight operating system where the time to compile is very low.
- Xv6 is a teaching operating system developed in the summer of 2006 for MIT

Developer	<u>MIT</u>
<u>Written in</u>	<u>C</u> and <u>assembly</u>
OS family	<u>Unix-like</u>
Source model	<u>Open source</u>

METHODOLOGY

Installation of QEMU

QEMU (short for Quick Emulator) is a hosted virtual machine monitor: it emulates CPUs through dynamic binary translation and provides a set of device models, enabling it to run a variety of unmodified guest operating systems. It also can be used with KVM to run virtual machines at near-native speed (requiring hardware virtualization extensions on x86 machines). QEMU can also do CPU emulation for user-level processes, allowing applications compiled for one architecture to run on another.

Run the following command to install qemu:

```
sudo apt-get install qemu
```

Installation of xv6

Xv6 is a re-implementation of the Unix sixth edition in order to use as a learning tool. xv6 was developed by MIT as a teaching operating system for their “6.828” course. A vital fact about xv6 is that it contains all the core Unix concepts and has a similar structure to Unix even though it lacks some functionality that you would expect from a modern operating system.

This is a lightweight operating system where the time to compile is very low and it also allow remote debugging.

Clone the official repository:

- **git clone <https://github.com/mit-pdos/xv6-public>**

Running xv6:

- **cd xv6**
- **make**
- **make qemu**

In order to define our own system call in xv6, we need to make changes to 5 files. Namely, these files are as follows.

Steps Followed To Add SYSTEM CALLS In LINUX KERNEL (XV6 OS):

- Created files named ps.c, clear.c, shutdown.c, nice.c, foo.c, cp.c within the folder xv6 containing all its files.
- Added the names and paths of all above mentioned files in Makefile.
- Added a new system call to the system call table in file syscall.h, usys.S, syscall.c.
- Added the relevant function call in sysproc.c.
- Added a piece of code that will determine what a system call will perform after it is called to proc.c.
- Finally we can run our system calls and see the outcome on the terminal screen.

1. syscall.h
2. syscall.c
3. sysproc.c
4. usys.S
5. user.h

syscall.h

```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat   8
10 #define SYS_chdir   9
11 #define SYS_dup    10
12 #define SYS_getpid  11
13 #define SYS_sbrk   12
14 #define SYS_sleep  13
15 #define SYS_uptime 14
16 #define SYS_open   15
17 #define SYS_write  16
18 #define SYS_mknod  17
19 #define SYS_unlink 18
20 #define SYS_link   19
21 #define SYS_mkdir  20
22 #define SYS_close  21
23 #define SYS_total   22
24 #define SYS_states  23
25 #define SYS_changeprS 24
```

syscall.c

```

85 extern int sys_chdir(void);
86 extern int sys_close(void);
87 extern int sys_dup(void);
88 extern int sys_exec(void);
89 extern int sys_exit(void);
90 extern int sys_fork(void);
91 extern int sys_fstat(void);
92 extern int sys_getpid(void);
93 extern int sys_kill(void);
94 extern int sys_link(void);
95 extern int sys_mkdir(void);
96 extern int sys_mknod(void);
97 extern int sys_open(void);
98 extern int sys_pipe(void);
99 extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_total(void);
107 extern int total_calls;
108 extern int sys_states(void);
109 extern int sys_changepr(void);
110

```

Usys.S

```

SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
SYSCALL(pipe)
SYSCALL(read)
SYSCALL(write)
SYSCALL(close)
SYSCALL(kill)
SYSCALL(exec)
SYSCALL(open)
SYSCALL(mknod)
SYSCALL(unlink)
SYSCALL(fstat)
SYSCALL(link)
SYSCALL(mkdir)
SYSCALL(chdir)
SYSCALL(dup)
SYSCALL(getpid)
SYSCALL(total)
SYSCALL(states)
SYSCALL changepr

```

User.h

```

4 // system calls
5 int fork(void);
6 int exit(void) __attribute__((noreturn));
7 int wait(void);
8 int pipe(int*);
9 int write(int, const void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int);
13 int exec(char*, char**);
14 int open(const char*, int);
15 int mknod(const char*, short, short);
16 int unlink(const char*);
17 int fstat(int fd, struct stat*);
18 int link(const char*, const char*);
19 int mkdir(const char*);
20 int chdir(const char*);
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
26 int total(void);
27 int states(void);
28 int changepr(int pid, int priority);

```

Default list of system calls in xv6

```

cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 n
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2170
cat        2 3 13436
echo       2 4 12508
forktest   2 5 8232
grep       2 6 15260
init       2 7 13100
kill       2 8 12552
ln         2 9 12460
ls         2 10 14676
mkdir     2 11 12572
rm         2 12 12556
sh         2 13 23196
stressfs   2 14 13228
usertests  2 15 56112
wc         2 16 14088
number     2 17 12744
zombie     2 18 12284
hello      2 19 12484
console    3 20 0
$ 

```

System calls that are added: -

1.total_sys - System Call which return the total number of method Calls in xv6 .

Sysproc.c

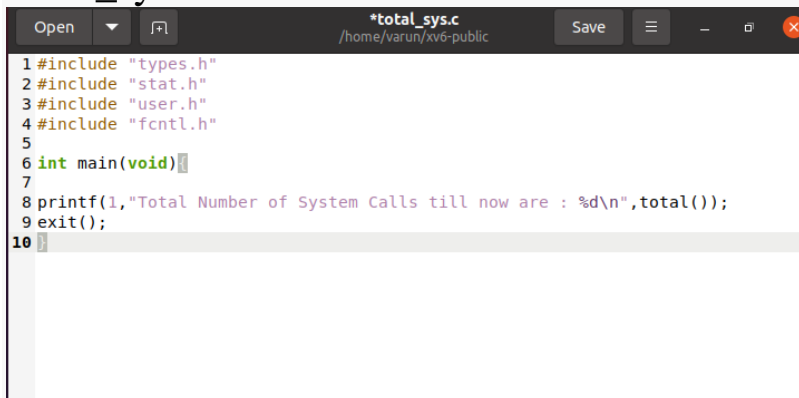
```
int sys_total(void)
{
    if(total_calls== -1) return total_calls;
    else return total_calls + 1;
}
```

Syscall.c

```
void
syscall(void)
{
    total_calls++;
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
    }
}
```

Total_sys.c



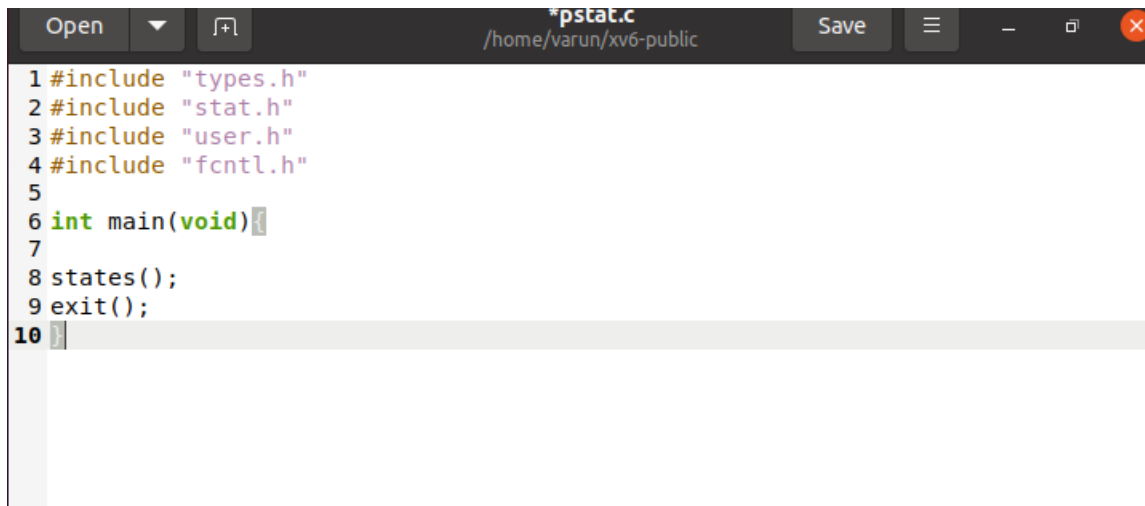
```
*total_sys.c
/home/varun/xv6-public

Open  Save  -  +  x

1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "fcntl.h"
5
6 int main(void)
7 {
8     printf(1, "Total Number of System Calls till now are : %d\n", total());
9     exit();
10 }
```


2.pstat - Shows the current status of process (Running , Sleeping , Runnable)

Pstat.c

A screenshot of a code editor window titled "pstat.c" with the path "/home/varun/xv6-public". The editor shows the first 10 lines of the C code. Lines 1-4 are include statements for "types.h", "stat.h", "user.h", and "fcntl.h". Line 5 is empty. Line 6 is the start of the main function. Lines 7-9 contain calls to "states()", "exit()", and an empty line. Line 10 is the start of a new block.

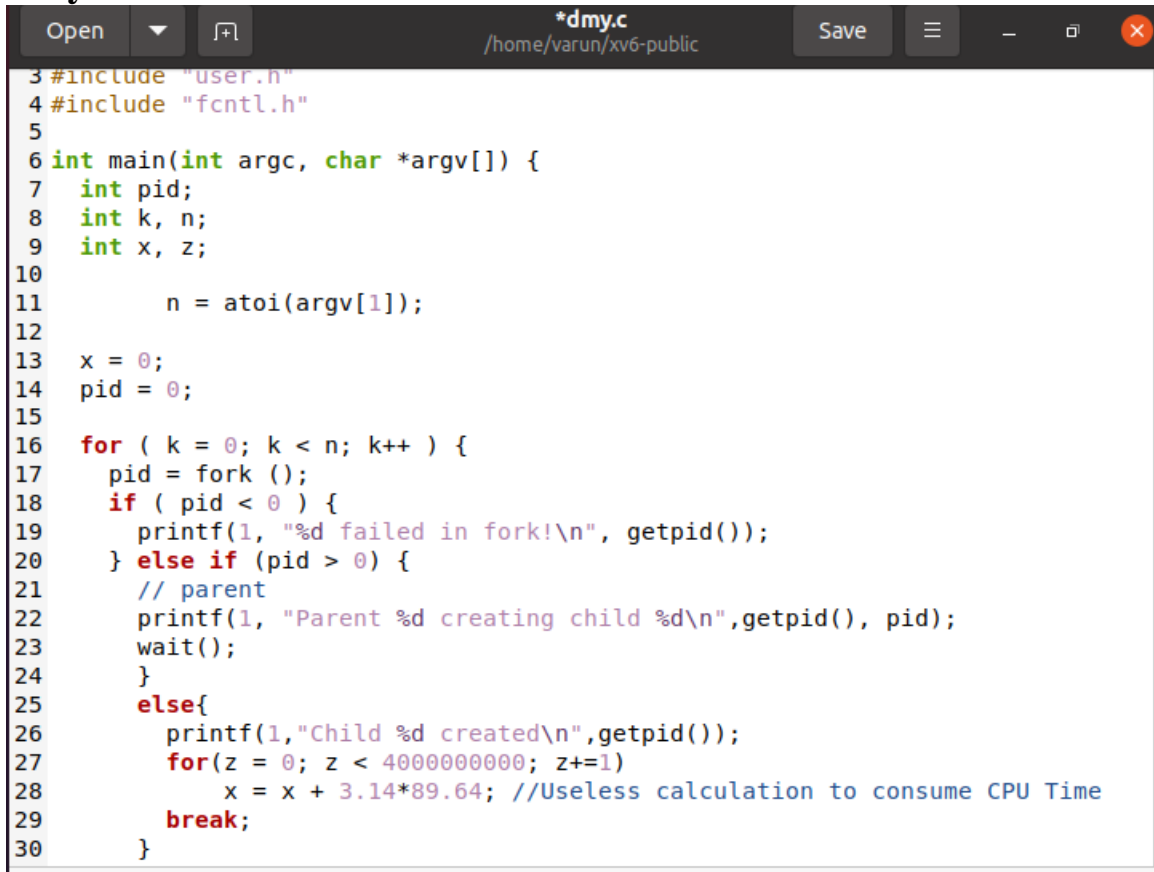
```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "fcntl.h"
5
6 int main(void){
7
8 states();
9 exit();
10
```

Proc.c

```
5 int states()
5 {
7
8 struct proc *p;
9
10 sti();
11
12
13
14 acquire(&ptable.lock);
15 cprintf("name \t pid \t state \t \tpriority\n");
16 for(p=ptable.proc; p<&ptable.proc[NPROC]; p++){
17 if(p->state==SLEEPING)
18 cprintf("%s \t %d \t SLEEPING \t %d \n" ,p->name,p->pid,p->priority);
19 if(p->state==RUNNING)
20 cprintf("%s \t %d \t RUNNING \t %d \n" ,p->name,p->pid,p->priority);
21 if(p->state==RUNNABLE)
22 cprintf("%s \t %d \t RUNNABLE \t %d \n" ,p->name,p->pid,p->priority);
23 }
24 release(&ptable.lock);
25
26 return 23;
27 }
```

3.dmy - used for dummy calculations by creating a parent-child processes.

Dmy.c



```
1 #include "user.h"
2 #include "fcntl.h"
3
4 int main(int argc, char *argv[]) {
5     int pid;
6     int k, n;
7     int x, z;
8
9     n = atoi(argv[1]);
10
11     x = 0;
12     pid = 0;
13
14     for ( k = 0; k < n; k++ ) {
15         pid = fork ();
16         if ( pid < 0 ) {
17             printf(1, "%d failed in fork!\n", getpid());
18         } else if ( pid > 0 ) {
19             // parent
20             printf(1, "Parent %d creating child %d\n",getpid(), pid);
21             wait();
22         } else{
23             printf(1,"Child %d created\n",getpid());
24             for(z = 0; z < 4000000000; z+=1)
25                 x = x + 3.14*89.64; //Useless calculation to consume CPU Time
26             break;
27         }
28     }
29 }
```

4.chgp - can be used to change the priority of the system call.

Proc.c

```
int changepr(int pid, int priority)
{
    struct proc *p;
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->pid == pid){
            p->priority = priority;
            break;
        }
    }
    release(&ptable.lock);
    return pid;
}
```

Sysproc.c

```
int sys_changepr(void){
    int pid,pr;
    if(argint(0,&pid)<0)
        return -1;

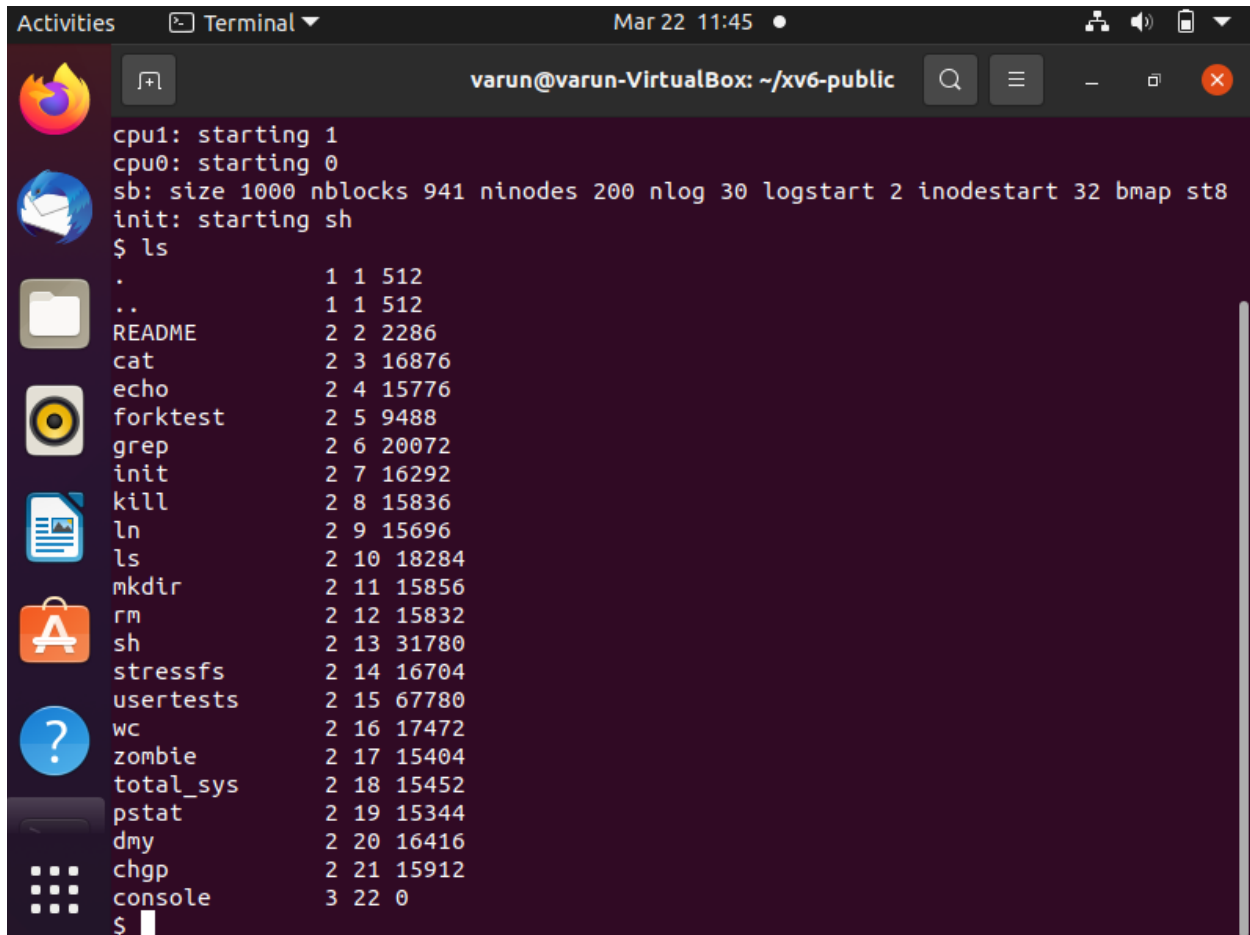
    if(argint(1,&pr)<0)
        return -1;

    return changepr(pid,pr);
}
```

Chgp.c

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "fcntl.h"
5
6 int main(int argc, char *argv[])
7 {
8     int priority, pid;
9
10    pid = atoi(argv[1]);
11    priority = atoi(argv[2]);
12    if (priority < 0 || priority > 20){
13        printf(2,"Invalid priority (0-20)!\n");
14        exit();
15    }
16
17    printf(1," pid=%d, pr=%d\n",pid,priority);
18    changepr(pid, priority);
19    exit();
20 }
```

OUTPUTS



```
Activities Terminal Mar 22 11:45
varun@varun-VirtualBox: ~/xv6-public

cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap st8
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat        2 3 16876
echo       2 4 15776
forktest  2 5 9488
grep       2 6 20072
init       2 7 16292
kill       2 8 15836
ln         2 9 15696
ls         2 10 18284
mkdir     2 11 15856
rm         2 12 15832
sh         2 13 31780
stressfs  2 14 16704
usertest  2 15 67780
wc         2 16 17472
zombie    2 17 15404
total_sys 2 18 15452
pstat     2 19 15344
dmy       2 20 16416
chgp      2 21 15912
console   3 22 0
$
```

Total_sys system call

```
$ total_sys
Total Number of System Calls till now are : 737
$ total_sys
Total Number of System Calls till now are : 803
$ total_sys
Total Number of System Calls till now are : 869
$
```

Dmy system call

```
$ dmy 3 &  
$ Parent 8 creating child 9  
Child 9 created  
dmy 3 8
```

Pstat system call

```
pstat  
name      pid      state      priority  
init      1        SLEEPING   10  
sh        2        SLEEPING   10  
dmy       9        RUNNABLE   10  
dmy       8        SLEEPING   10  
dmy       12       RUNNING    10  
dmy       11       SLEEPING   10  
pstat     13       RUNNING    10  
$
```

Chgp system call

```

$ chgp 8 6
  pid=8, pr=6
$ pstat
name      pid      state      priority
init       1      SLEEPING    10
sh         2      SLEEPING    10
dmy        9      RUNNING     10
dmy        8      SLEEPING     6
dmy       12      RUNNABLE    10
dmy       11      SLEEPING    10
pstat     15      RUNNING     10
$ chgp 11 4
  pid=11, pr=4
$ pstat
name      pid      state      priority
init       1      SLEEPING    10
sh         2      SLEEPING    10
dmy        9      RUNNABLE    10
dmy        8      SLEEPING     6
dmy       12      RUNNING     10
dmy       11      SLEEPING     4
pstat     17      RUNNING     10
$ s

```

```

$ pstat
name      pid      state      priority
init       1      SLEEPING    10
sh         2      SLEEPING    10
dmy        9      RUNNING     8
dmy        8      SLEEPING     6
dmy       12      RUNNABLE    10
dmy       11      SLEEPING     4
pstat     20      RUNNING     10

```

CONCLUSION

- We can add system calls such as **create child processes** with our own fork system call and also control its **run time** before it becomes zombie.
- We deep dive into the very basics of an Operating system and make changes in it at root level. We can add features to any Open Source OS similarly.
- Some system calls do not take any arguments and return just an integer value (e.g., uptime in sysproc.c). Some other system calls take in multiple arguments like strings and integers (e.g., open system call in sysfile.c), and return a simple integer value. Further, more complex system calls return a lot of information back to the user program in a user-defined structure.

REFERENCES

- <https://medium.com/@silvamatteus/adding-user-programs-to-xv6-ba9896605942>
- <https://en.wikipedia.org/wiki/Xv6>
- <https://github.com/mit-pdos/xv6-public>
- <https://www.ijsr.net/archive/v6i1/5011702.pdf>